

Concurrent Queues

George Poulos
gmp289@cs.nyu.edu

Simon Tso
st1450@nyu.edu

Introduction

The need for software concurrency has grown rapidly over the last decade. As raw processor speed has steadied over the last few years, the number of cores per processor has increased. This has further emphasized the importance of concurrent software. However, introducing concurrency into most programming models has not been effortless. There is a great deal of complexity involved in properly harnessing the power of multicore systems. One aspect of this complexity can be seen in concurrent data structures. Specifically, in this report will describe our design and report on our implementations of concurrent queues.

Queues in General

Queues are a elementary data structure. It is a container with first-in-first-out semantics. A Queue interface consists of two basic operations for adding and removing items into the queue. Here is the interface we used for our queues.

```
interface IQueue<T> {  
    void Enqueue(T value);  
    bool Dequeue(T *value);  
}
```

Enqueue will add the value into the queue. Dequeue will attempt to remove a value from the queue. If the queue is empty, it will return false. Otherwise, it will fill in the supplied pointer with the dequeued value and return true.

A simple implementation of a queue uses a list of nodes and has a pointer to the head and the tail. When an enqueue is performed, the item is added to the back of the list and tail pointer is adjusted. Similarly, when a dequeue is performed the first item is removed from the list and the head pointer is adjusted. This is where the complexity in implementing concurrent queues begins. There is a need to perform two actions at once; to link/unlink the item and to adjust the head/tail pointers. We explored two implementations that allow for this behavior to occur concurrently.

As a note, a key factor that allows both our queue implementations to work properly is the idea of using a sentinel node inside the queue. Initially, the queue will start with a dummy node that allows each operation to blindly follow Next pointers without checking for null. Additionally, this allows each operation to only need either the head or tail pointer and not both.

Here is pseudo code for our simple queue implementation.

```
SimpleQueue<T>() {
    Node<T> Node = new Node<T>();
    node->Next = NULL;
    head = tail = node;
}

void Enqueue(T value) {
    Node<T>* node = new Node<T>();
    node->Value = value;
    node->Next = NULL;
    tail->Next = node;
    tail = node;
}

bool Dequeue(T* value) {
    Node<T>* node = head;
    Node<T>* next = node->Next;
    if(!next) return false;
    *value = next->Value;
    head = next;
    delete node;
    return true;
}
```

Locking Queues

Using a lock has been the most straight-forward and popular implementation in the past. Locks are somewhat simple compared to other implementations. A locking queue allows enqueues and dequeues to occur concurrently and safely. The downside to the lock is lock contention. Since a lock only grants access of each operation to one thread at a time, a large performance hit can occur when there are many threads trying to acquire the lock. Lock acquisition actually serializes each thread that is trying to perform an operation that another thread is also trying to perform.

Here is an example of adding a lock to the Enqueue operation for a locking queue.

```
void Enqueue(T value) {
    Node<T>* node = new Node<T>();
    node->Value = value;
    node->Next = NULL;
    lock(EnqueueMutex){
        tail->Next = node;
        tail = node;
    }
}
```

Lockless Queues

One way to avoid lock acquisition is not to use locks at all. An alternative to locks can be atomic operations. Using atomic operations requires a somewhat more complicated thought process. One powerful atomic operation is compare and swap (CAS). Using CAS, a lockless queue can be implemented. However, CAS is not a silver bullet. CAS has a few disadvantages up front. First, heavy use of CAS between threads will invalidate CPU caches. This can cause coherency traffic between the cores and tremendously decrease performance. Secondly, the data structure may fall into intermediate states during operations. It is therefore required to handle all intermediate states during the course of

normal operation. Despite these shortcomings, it is possible to comply with these caveats and implement a concurrent queue using CAS.

In general, CAS operations work by taking snapshots of memory. Values in memory that are required to perform the operation will be copied locally. Then the execution will enter a loop and prepare its local structure to be swapped into the queue structure. A CAS will be performed to check if the snapshot has changed and attempt to swap in the prepared value. If the snapshot is unchanged, the swap will succeed and the execution will break out of the loop. If the snapshot has changed, the execution will loop and try again by taking a new snapshot and continuing from there.

Here is an implementation of enqueue using CAS instead of locks.

```
void Enqueue(T value) {
    Node<T>* node = new Node<T>();
    node->Value = value;
    node->Next = NULL;

    Node<T>* t;
    Node<T>* next;
    while(true){
        t = Tail;
        if(Tail != t) continue;
        next = t->Next;
        if(Tail != t) continue;
        if(next){ CAS(&Tail, t, next); continue; }
        if(CAS(&t->Next, NULL, node)) break;
    }
    CAS(&Tail, t, node);
}
```

Further Problems with CAS

With proper consideration, avoiding to use CAS recklessly and handling intermediate states is fairly easy. However, there are two addition problems that arise when using CAS in a queue. These problems can be subtle at first glance and while they are related, they are indeed separate problems.

The first problem is the memory reclamation problem. This problem is caused by the inability to detect when a node is not use by any threads. Usually, in a dequeue operation, the calling thread deletes the node that was removed the queue. There is a possibility though that other threads are accessing this node in a CAS operation. Therefore, it is unsafe to delete the node at the time of a dequeue. Determining when it is safe to free memory of the nodes in the queue is the memory reclamation problem.

The second problem is the ABA problem. The ABA problem occurs when after a thread takes a snapshot for a CAS operation, another thread deletes a node in the snapshot, frees it and allocates a new node that just happens to have the same memory address of the original. This will allow the

snapshot to appear unchanged even though the node is really a new one. Avoiding this scenario will solve the ABA problem.

Hazard Pointers

Since both problems are related, they can be solve collectively through a mechanism know as Hazard Pointers. Hazard Pointers provide a way for threads to broadcast which nodes are currently being used in their CAS snapshots. This is referred to as declaring pointers as hazards. Upon completion of a dequeue, the freeing of the dequeued node will be deferred to a later time instead of being freed immediately. This is known as retiring the node. Once some preset amount of nodes have been retired, a scan is done and any nodes not declared hazards by any threads can be safely freed.

To accomplish this, a structure called a Hazard Pointer Record (HPRec) is used. The queue maintains an internal chain of HPRecs. When a thread requires access to the queue, it is given an HPRec.

```
struct HPRec {
    Node* HP[K];
    HPRec* Next;
    bool Active;
    list<Node*> RetireList;
}
```

Inside the HPRec is a an array of Hazard Pointers (HP[]). When a thread takes a snapshot for a CAS and needs to declare some pointers as hazards, it inserts them into this array. Each thread can only modify the HP[] array inside its HPRec but can read all others in the queue. The Next pointer is used by the queue to maintain the chain of HPRecs. The Active boolean field is used to deactivate a HPRec when a thread finishes and save it for reuse when a new thread asks for access to the queue. When a dequeue occurs, the node is retired into the RetireList of the calling thread.

Since an HPRec is given to each thread, each thread has its own RetireList. Once this list grows large enough, a scan is performed. A scan is similar to a garbage collection where it will determine which nodes can be freed based on whether they are in use by any threads. A scan consists of two parts.

```
void scan(){
    map<Node*,bool> hazards;
    foreach(HPRec rec in HPRecChain)
        foreach(Node* hp in rec.HP)
            hazards.insert(hp);

    foreach(Node* node in RetireList)
        if(!hazards.lookup(node)){
            RetireList.remove(node);
            delete node;
        }
}
```

The first part of the scan will traverse the queue's HPRec chain and store any hazards into a local map. This map will be used in part 2 of the scan to quickly lookup pointers and see if they are hazards. The map is used to avoid the need to traverse the entire HPRec chain to do a lookup.

The second part of the scan will then iterate through the local RetireList of the calling thread. For each node in the RetireList, it will perform a lookup in the map created in part 1. If the node is found in the map, it has been declared a hazard by another thread and is unsafe to delete. As a result, it will simply be left in the RetireList. On the other hand, if the node is not found in the map, it is not in use by any other threads and safely freed. It is then removed from the RetireList and deleted.

The next part of using Hazard Pointers is determining what the hazards are. Most data structures will only require one or two hazards per thread at a time. A hazard analysis must be done to determine exactly how many are needed. Hazard pointers are by definition the pointers that are unsafe to delete at a given time. The way pointers become unsafe is when they are referenced by one or more threads. The way this usually happens is during a CAS snapshot. Looking at the points that are used during these snapshots will determine the number of hazard pointers are needed for the entire structure.

In the case of the queue, since an enqueue and a dequeue from the same thread cannot be going on at the same time, the number of hazards needed per thread is equal to the max of the ones used during an enqueue or a dequeue. During an enqueue operation, a snapshot is only utilizing the tail pointer. A dequeue operation however takes a snapshot using both the head and the head->next pointers. Therefore, the size of the HP[] inside a HPRec is 2. Once this analysis is done, it is simple just a matter of storing the pointers into the HP[] at the right time as seen here.

```
void Enqueue(T value) {
    Node<T>* node = new Node<T>();
    node->Value = value;
    node->Next = NULL;

    Node<T>* t;
    Node<T>* next;
    while(true){
        t = Tail;
        hprec->HP[0] = t;
        if(Tail != t) continue;
        next = t->Next;
        if(Tail != t) continue;
        if(next){ CAS(&Tail, t, next); continue; }
        if(CAS(&t->Next, NULL, node)) break;
    }
    CAS(&Tail, t, node);
}
```

Using the LocklessQueue Implementation

As described earlier, we created an `IQueue<T>` interface that is implemented directly by our `SimpleQueue` and `LockingQueue`. However, the `LocklessQueue` is unable to do this because each thread is required to have an `HPRec` and be able to access it somehow. There are a few ways to go about accomplishing this. Thread Local Storage could be used to save a reference to an `HPRec` and then that reference can be passed to an `Enqueue(HPRec* rec, T value)` method. However, this breaks the abstraction of using an `IQueue<T>` in first place. The way we implemented a solution was to use something we call a `ThreadAccessor`.

A `ThreadAccessor` is a private class of `LocklessQueue` that implements the `IQueue<T>` interface. When a thread needs access to a queue, it calls the `CreateAccessor()` method of an instance of `LocklessQueue`. This will return an `IQueue<T>` reference that the thread can use to access the queue. The `IQueue<T>` reference in actuality is an instance of `ThreadAccessor` that encapsulates the `HPRec` for that particular thread. A slight subtlety to make this work correctly was declaring `ThreadAccessor` a friend of `LocklessQueue`. Friend classes can break encapsulation but since `ThreadAccessor` is a private class of `LocklessQueue` it seemed like a legitimate solution. As an example, here is how the `LocklessQueue` can be used.

```
LocklessQueue<int> lockless;
IQueue<int>* a = lockless.CreateAccessor();
int x = 0;
a->Enqueue(1);
a->Enqueue(2);
a->Dequeue(&x); // returns true, x = 1
delete a;
```

Performance

We have performed a few performance tests using both the locking and lockless queues. The system we used for testing was energon (8 core 1.8GHz Intel Xeon). Our tests have evolved over the course of the project. In this report, we will present the tests based on the bench program found in bench.cpp.

This program runs two types of tests. The first one runs either a enqueue or dequeue randomly. The second runs a series of enqueues followed by a series of dequeues. The series test runs twice, once for a bias towards and enqueues and then with a bias towards dequeues.

In addition to the concurrent tests, a sequential test is also done. A simple queue implementation is added which is identical to the locking queue without the locks. This is to compare the single-threaded performance of a naive queue versus the concurrent queue implementations.

To simulate a workload, the benchmark can optionally run a sieve function in between queue operations. A sieve function calculates the primes up to an given upper bound. Increasing the upper bound will cause the sieve to run longer causing the queue operations to be more spread out over time

All the tests are run using integer typed queues. The tests also run a correctness check, which accumulates the sum of enqueued items. Another sum is used for dequeued items. These sums are compared and adjusted (when the queue does not start or finish empty). If they are equal at the end of a test, it can be assumed that the queue has not failed in correctness.

Compiling

To compile, we used the following command:

```
g++ IQueue.h SimpleQueue.h LockingQueue.h LocklessQueue.h bench.cpp  
-Wall -lrt -lpthread -o bench
```

-lrt is used to take timestamps and only required for the bench.cpp test. If just compiling the queues, the only library you need to include is pthread.

Test Results

Time shown is displayed in the last column in milliseconds.

Iterations 1000000
Sieve Bound 0
Threads 2
Series Bias 2

Random Tests			
Sequential Simple	PASS	62634	
Sequential Locking	PASS	125316	
Sequential Lockless	PASS	465140	
Concurrent Locking	PASS	680483	
Concurrent Lockless	PASS	687245	

Enqueue Bias Series Tests			
Sequential Simple	PASS	120932	
Sequential Locking	PASS	200324	
Sequential Lockless	PASS	469473	
Concurrent Locking	PASS	994875	
Concurrent Lockless	PASS	571673	

Dequeue Bias Series Tests			
Sequential Simple	PASS	147297	
Sequential Locking	PASS	236097	
Sequential Lockless	PASS	823506	
Concurrent Locking	PASS	1359808	
Concurrent Lockless	PASS	896773	

Iterations 1000000
Sieve Bound 0
Threads 4
Series Bias 2

Random Tests			
Sequential Simple	PASS	61954	
Sequential Locking	PASS	125904	
Sequential Lockless	PASS	465275	
Concurrent Locking	PASS	1369490	
Concurrent Lockless	PASS	488774	

Enqueue Bias Series Tests			
Sequential Simple	PASS	122865	
Sequential Locking	PASS	199504	
Sequential Lockless	PASS	469541	
Concurrent Locking	PASS	1559039	
Concurrent Lockless	PASS	434948	

Dequeue Bias Series Tests			
Sequential Simple	PASS	137535	
Sequential Locking	PASS	243589	
Sequential Lockless	PASS	816693	
Concurrent Locking	PASS	1815145	
Concurrent Lockless	PASS	611822	

Iterations 1000000
Sieve Bound 0
Threads 8
Series Bias 2

Random Tests			
Sequential Simple	PASS	62666	
Sequential Locking	PASS	124256	
Sequential Lockless	PASS	467526	
Concurrent Locking	PASS	1020029	
Concurrent Lockless	PASS	425888	

Enqueue Bias Series Tests			
Sequential Simple	PASS	129104	
Sequential Locking	PASS	202686	
Sequential Lockless	PASS	471317	
Concurrent Locking	PASS	863716	
Concurrent Lockless	PASS	529741	

Dequeue Bias Series Tests			
Sequential Simple	PASS	133589	
Sequential Locking	PASS	237955	
Sequential Lockless	PASS	838871	
Concurrent Locking	PASS	913842	
Concurrent Lockless	PASS	496769	

Iterations 1000000
Sieve Bound 0
Threads 16
Series Bias 2

Random Tests			
Sequential Simple	PASS	61470	
Sequential Locking	PASS	124619	
Sequential Lockless	PASS	466233	
Concurrent Locking	PASS	738649	
Concurrent Lockless	PASS	499883	

Enqueue Bias Series Tests			
Sequential Simple	PASS	126330	
Sequential Locking	PASS	200271	
Sequential Lockless	PASS	441234	
Concurrent Locking	PASS	885405	
Concurrent Lockless	PASS	500252	

Dequeue Bias Series Tests			
Sequential Simple	PASS	135313	
Sequential Locking	PASS	238540	
Sequential Lockless	PASS	840656	
Concurrent Locking	PASS	980647	
Concurrent Lockless	PASS	669595	

Iterations 1000000
Sieve Bound 100
Threads 2
Series Bias 2

Random Tests			
Sequential Simple	PASS		1684332
Sequential Locking	PASS		1743544
Sequential Lockless	PASS		2158061
Concurrent Locking	PASS		1676784
Concurrent Lockless	PASS		1631442

Enqueue Bias Series Tests			
Sequential Simple	PASS		1838472
Sequential Locking	PASS		1917211
Sequential Lockless	PASS		2206580
Concurrent Locking	PASS		1362624
Concurrent Lockless	PASS		1620232

Dequeue Bias Series Tests			
Sequential Simple	PASS		1954931
Sequential Locking	PASS		2061635
Sequential Lockless	PASS		3005915
Concurrent Locking	PASS		1563148
Concurrent Lockless	PASS		1891935

Iterations 1000000
Sieve Bound 100
Threads 4
Series Bias 2

Random Tests			
Sequential Simple	PASS		1685974
Sequential Locking	PASS		1738497
Sequential Lockless	PASS		2133310
Concurrent Locking	PASS		1805870
Concurrent Lockless	PASS		918474

Enqueue Bias Series Tests			
Sequential Simple	PASS		1840802
Sequential Locking	PASS		1913413
Sequential Lockless	PASS		2207364
Concurrent Locking	PASS		1541419
Concurrent Lockless	PASS		836019

Dequeue Bias Series Tests			
Sequential Simple	PASS		2286186
Sequential Locking	PASS		2392008
Sequential Lockless	PASS		2723637
Concurrent Locking	PASS		1357968
Concurrent Lockless	PASS		1072946

Iterations 1000000
Sieve Bound 100
Threads 8
Series Bias 2

Random Tests			
Sequential Simple	PASS		1681155
Sequential Locking	PASS		1733117
Sequential Lockless	PASS		2124352
Concurrent Locking	PASS		1118121
Concurrent Lockless	PASS		624164

Enqueue Bias Series Tests			
Sequential Simple	PASS		1828144
Sequential Locking	PASS		1918603
Sequential Lockless	PASS		2202487
Concurrent Locking	PASS		1076796
Concurrent Lockless	PASS		585136

Dequeue Bias Series Tests			
Sequential Simple	PASS		1867307
Sequential Locking	PASS		2041704
Sequential Lockless	PASS		2712062
Concurrent Locking	PASS		1026498
Concurrent Lockless	PASS		706690

Iterations 1000000
Sieve Bound 100
Threads 16
Series Bias 2

Random Tests			
Sequential Simple	PASS		1687690
Sequential Locking	PASS		1735187
Sequential Lockless	PASS		2196187
Concurrent Locking	PASS		832173
Concurrent Lockless	PASS		711443

Enqueue Bias Series Tests			
Sequential Simple	PASS		1819161
Sequential Locking	PASS		1914556
Sequential Lockless	PASS		2205970
Concurrent Locking	PASS		1133394
Concurrent Lockless	PASS		625436

Dequeue Bias Series Tests			
Sequential Simple	PASS		1858490
Sequential Locking	PASS		2049763
Sequential Lockless	PASS		2728057
Concurrent Locking	PASS		1216327
Concurrent Lockless	PASS		829559

Iterations 100000
Sieve Bound 1000
Threads 16
Series Bias 2

Random Tests			
Sequential Simple	PASS	1710704	
Sequential Locking	PASS	1714754	
Sequential Lockless	PASS	1750059	
Concurrent Locking	PASS	270345	
Concurrent Lockless	PASS	297413	

Enqueue Bias Series Tests			
Sequential Simple	PASS	1722922	
Sequential Locking	PASS	1721790	
Sequential Lockless	PASS	1753386	
Concurrent Locking	PASS	284296	
Concurrent Lockless	PASS	302928	

Dequeue Bias Series Tests			
Sequential Simple	PASS	1724600	
Sequential Locking	PASS	1733508	
Sequential Lockless	PASS	1776346	
Concurrent Locking	PASS	269701	
Concurrent Lockless	PASS	315333	

Iterations 100000
Sieve Bound 1000
Threads 96
Series Bias 2

Random Tests			
Sequential Simple	PASS	1711124	
Sequential Locking	PASS	1716284	
Sequential Lockless	PASS	1744635	
Concurrent Locking	PASS	250760	
Concurrent Lockless	PASS	323766	

Enqueue Bias Series Tests			
Sequential Simple	PASS	1726977	
Sequential Locking	PASS	1730159	
Sequential Lockless	PASS	1754679	
Concurrent Locking	PASS	243276	
Concurrent Lockless	PASS	261952	

Dequeue Bias Series Tests			
Sequential Simple	PASS	1729141	
Sequential Locking	PASS	1732679	
Sequential Lockless	PASS	1776111	
Concurrent Locking	PASS	250506	
Concurrent Lockless	PASS	261257	

Iterations 100000
Sieve Bound 100
Threads 96
Series Bias 2

Random Tests			
Sequential Simple	PASS	169286	
Sequential Locking	PASS	173489	
Sequential Lockless	PASS	212566	
Concurrent Locking	PASS	93006	
Concurrent Lockless	PASS	129832	

Enqueue Bias Series Tests			
Sequential Simple	PASS	181989	
Sequential Locking	PASS	186203	
Sequential Lockless	PASS	215674	
Concurrent Locking	PASS	113348	
Concurrent Lockless	PASS	65567	

Dequeue Bias Series Tests			
Sequential Simple	PASS	178142	
Sequential Locking	PASS	184810	
Sequential Lockless	PASS	243683	
Concurrent Locking	PASS	116422	
Concurrent Lockless	PASS	99960	

Iterations 100000
Sieve Bound 0
Threads 96
Series Bias 2

Random Tests			
Sequential Simple	PASS	6021	
Sequential Locking	PASS	12309	
Sequential Lockless	PASS	46734	
Concurrent Locking	PASS	89177	
Concurrent Lockless	PASS	84244	

Enqueue Bias Series Tests			
Sequential Simple	PASS	12562	
Sequential Locking	PASS	17438	
Sequential Lockless	PASS	42747	
Concurrent Locking	PASS	91847	
Concurrent Lockless	PASS	56740	

Dequeue Bias Series Tests			
Sequential Simple	PASS	10005	
Sequential Locking	PASS	16251	
Sequential Lockless	PASS	66385	
Concurrent Locking	PASS	73489	
Concurrent Lockless	PASS	59347	

The test results are promising. First off, in a sequential scenario, the concurrent locks get destroyed in performance. The overhead introduced to handle the concurrency is a real burden when it is not being used. The first four results show raw queue access with no sieve being performed. There is literally no motivation here to use a concurrent queue. This is not a real-world scenario though so comparing the concurrent queues to sequential queues is of little value. It is interesting to note though that with very high concurrent access, the lockless queue is much faster than the locking queue.

Moving onto the next set of four, these show a small sieve and are differentiated by number of threads. In the two-thread case, the overhead of the concurrency is immediately diminished and all queues are about the same except for the lockless which is far behind. As the number of threads increases, the concurrent queues pull farther ahead. Also, as the lock contention starts to grow, the lockless threads sees some even further gains.

The next test result shows a large sieve with 16 threads. The concurrent queues performance is magnitudes ahead of the sequential here as the execution is CPU bound. Here we can see that the lockless queue has lost its glory and has fallen behind the locking queue slightly. Since the majority of time is spent in computation, the contention for the lock has become a smaller part of the execution.

The last three results show a small sieve with a large, 96, amount of threads. The results follow the trend of the more lock contention there is, the more benefit of the lockless queue.

Notes and Overall Remarks

There are few things that should be noted in regard to the overall project. First, there are few things that could be tuned to increase performance. The frequency of the scans can be adjusted to increase performance at the cost of using more memory. Also, the list and maps used during the scan can most likely be optimized. In addition, pooling nodes can drastically reduce the times for enqueues. If memory is pre-allocated, the intuition is that there enqueueing will be much faster.

Testing also turned out to be more difficult than first conceived. Our testing tools have evolved over the course of the project and even now we are not sure they are completely reliable. A lot of performance gains by using Hazard Pointers can be machine and usage dependent. It is a very hard thing to test and feel completely confident about. A more purposeful test would be to swap in the lockless queue into a system that is already using a concurrent queue, tune it correctly and then observe the results compared to the original.

As an additional note, there is a slightly subtlety in the HPRC chain mechanism. The queue is unable to actually delete any nodes from the HPRC chain. Trying to do so using CAS would be a catch 22 since safely reclaiming memory is the problem the HPRC chain is used to solve. The algorithm simply deactivates nodes instead of deleting them. Any deactivated nodes are reused when possible instead of allocating new ones. This avoids the problem completely.

Conclusion

Using a lockless queue is not an automatic solution to any concurrency woes. It should be observed as a tool to help fine-tune performance where it can be used appropriately. The simple implementation and usage of the locking queue can sometimes outweigh the lockless queue in both complexities and performance. It is important to note though that both the locking and lockless queues can embrace the power of multicore systems. It is up to the programmer to correctly identify and use the proper techniques for solving every problem.

References

Our implementation is primarily based on the first paper listed by Maged Michael.

Hazard Pointers: Safe Memory Reclamation for Lock-Free Objects

Maged M. Michael

<http://www.research.ibm.com/people/m/michael/ieeetpds-2004.pdf>

Simple, Fast, and Practical Non-Blocking and Blocking Concurrent Queue Algorithms

Maged M. Michael; Michael L. Scott

http://www.cs.rochester.edu/~scott/papers/1996_PODC_queues.pdf

The Art of Multiprocessor Programming

Herlihy, Maurice, and Nir Shavit. Morgan Kaufmann, 2008. Print.

CAS-Based Lock-Free Algorithm for Shared Deques

Maged M. Michael

<http://www.research.ibm.com/people/m/michael/europar-2003.pdf>

Lock-free Dynamically Resizable Arrays

Damian Dechev; Peter Pirkelbauer; Bjarne Stroustrup

<http://www2.research.att.com/~bs/lock-free-vector.pdf>